

Music Generation using WaveNet Architecture

Karthik Nagarajan Sundar

Chalmers University of Technology
Electrical Engineering Department



CHALMERS
UNIVERSITY OF TECHNOLOGY

Introduction

Main objective: Prediction of subsequent note/chord for a certain period of time to generate a track of music using Deep learning.

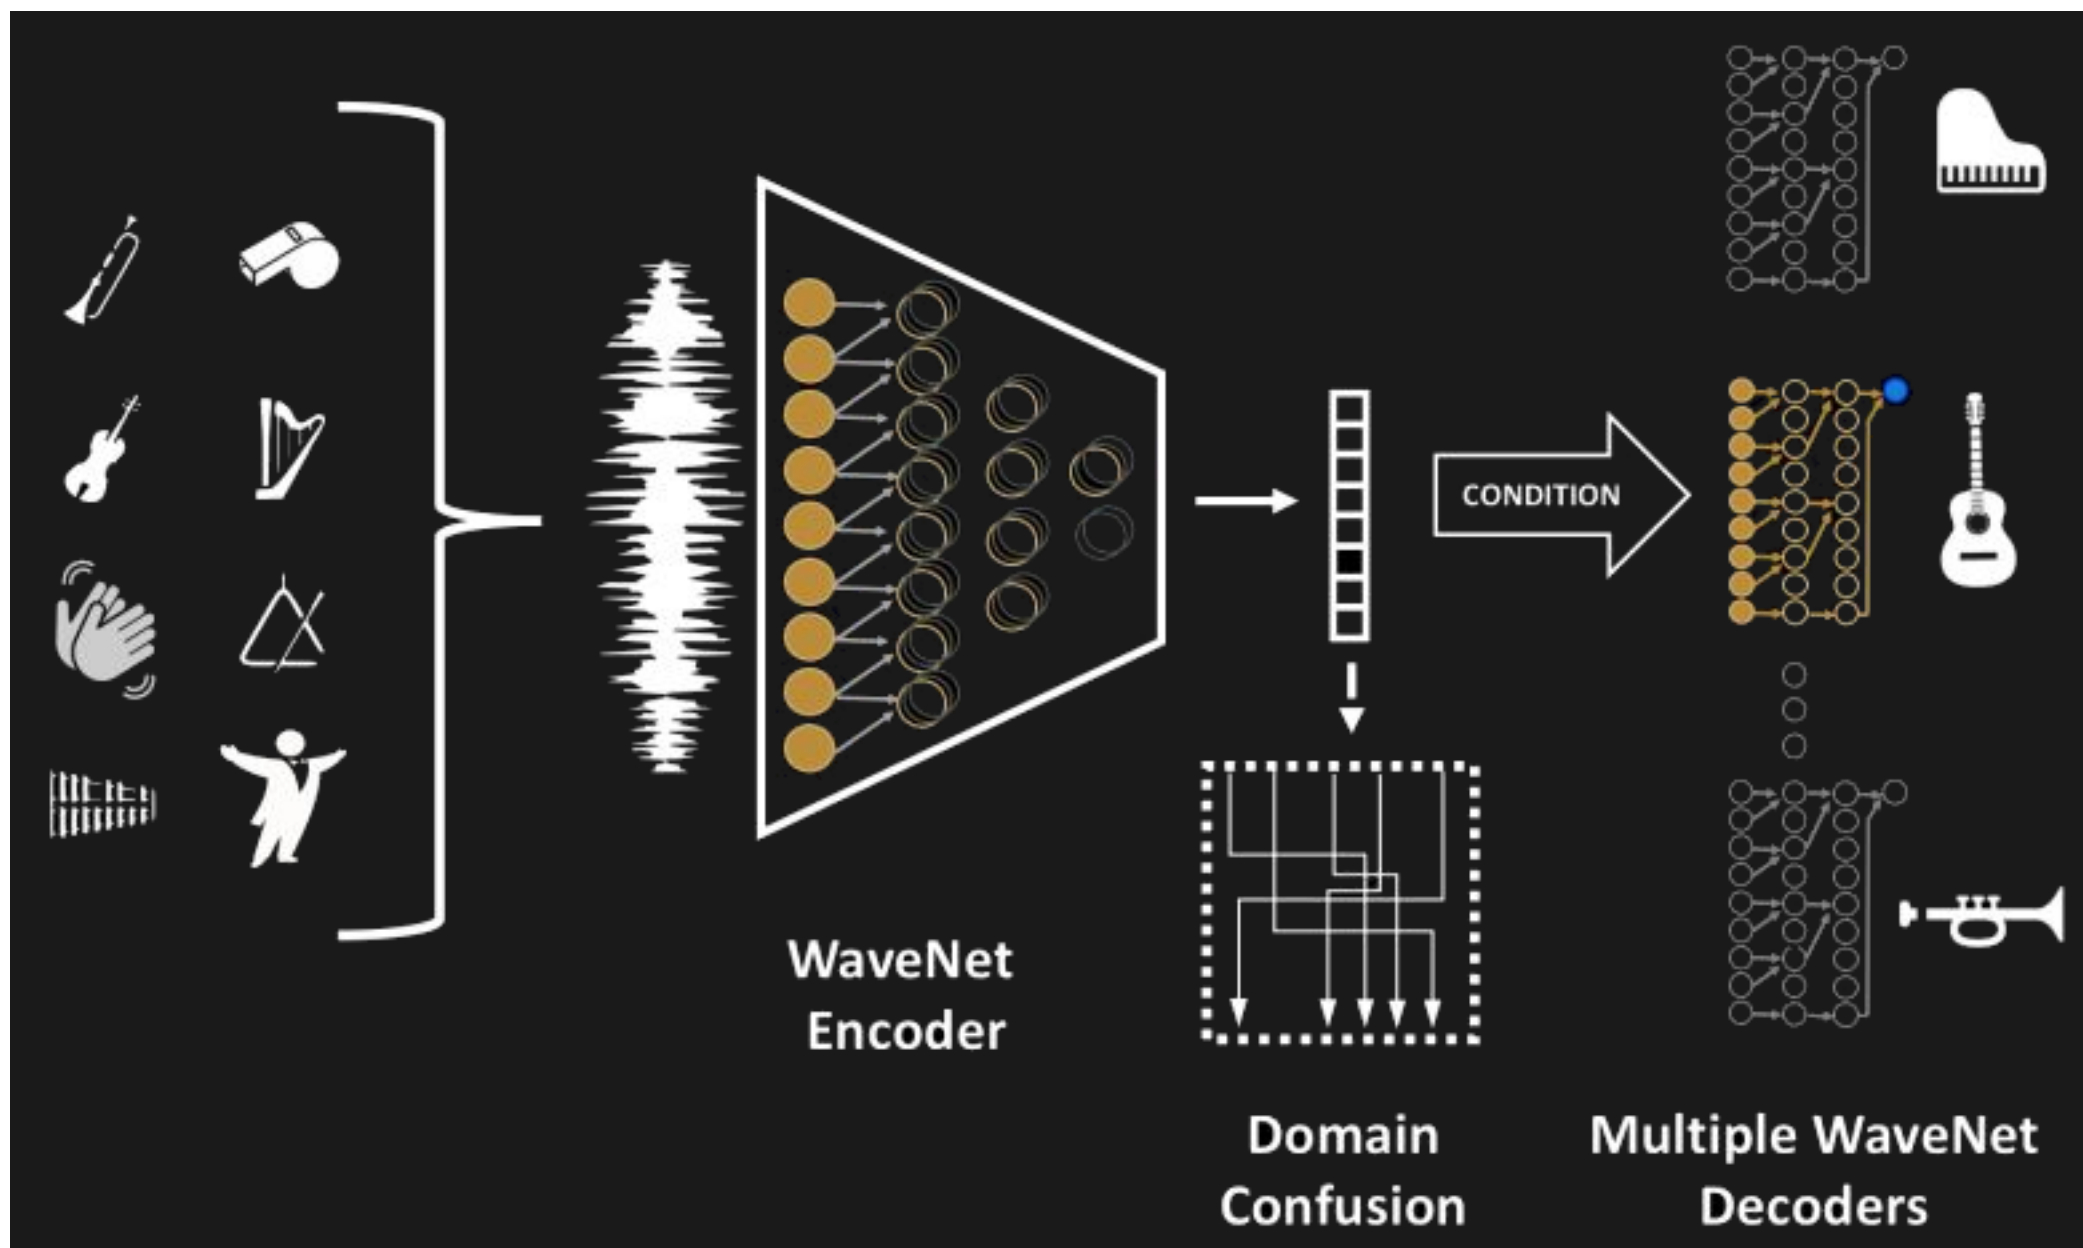


Figure 1: Music generation - basic idea

Contents:

- ▶ Music dataset(midi format)
 - ▶ Reading a midi file
 - ▶ Visualizing midi file in terms of Piano rolls.
- ▶ WaveNet architecture
 - ▶ What is Wavenet
 - ▶ Model summary
- ▶ Results and Evaluation

Music Data set

- ▶ Self-compiled famous metal genre music (Total:65)
- ▶ Format used for processing: Musical Instrument Digital Interface (MIDI)
- ▶ Split ratio: 80/20 Training and Validation data-sets
- ▶ Reading a midi file
 - ▶ Notes: sound produced by a single key
 - ▶ Chord: sound produced by 2 or more keys simultaneously




Figure 2: Visualizing midi in terms of piano rolls

WaveNet

- ▶ Deep Learning-based generative model developed by DeepMind(Google)
- ▶ Input: a sequence of nodes and chords
- ▶ Output: new predicted samples from the sequence
- ▶ uses Dilated Convolution layer instead of Causal Convolution layers
- ▶ Major difference: all the previous inputs can have significant contribution on predicted notes




Figure 3: Diluted Convolution Layer

- ▶ Better than RNNs because
 - ▶ better long term memory
 - ▶ No vanishing gradients problem
 - ▶ Hardware friendly
- ▶ Better than LSTMs because
 - ▶ longer memory for sequences over 10k samples.

Model Summary

- ▶ 3 1-dimensional Convolution layers
- ▶ Dropout layers: avoid over fitting
- ▶ Max-pooling layers: Down sample

Model: "sequential"		
Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 32, 100)	23300
conv1d (Conv1D)	(None, 32, 512)	154112
dropout (Dropout)	(None, 32, 512)	0
max_pooling1d (MaxPooling1D)	(None, 16, 512)	0
conv1d_1 (Conv1D)	(None, 16, 128)	196736
dropout_1 (Dropout)	(None, 16, 128)	0
max_pooling1d_1 (MaxPooling1D)	(None, 8, 128)	0
conv1d_2 (Conv1D)	(None, 8, 16)	6160
dropout_2 (Dropout)	(None, 8, 16)	0
max_pooling1d_2 (MaxPooling1D)	(None, 4, 16)	0
global_max_pooling1d (GlobalMaxPooling1D)	(None, 16)	0
dense (Dense)	(None, 256)	4352
dense_1 (Dense)	(None, 235)	60395
Total params: 445,055		
Trainable params: 445,055		
Non-trainable params: 0		

Figure 4: Model Summary

Results and Evaluation

- ▶ Parameters: 75 epochs run and batch size of 128
- ▶ Number of unique and random notes generated: 17 for a sequence length of 30 notes/chords
- ▶ User feedback: 5 out of 7 volunteers liked the track




Figure 5: Accuracy




Figure 6: Loss




Figure 7: Visualizing output in terms of Piano rolls

References

- [1] Aravind Pai, "Automatic music generation", 2020, <https://www.analyticsvidhya.com/blog/2020/01/how-to-perform-automatic-music-generation/>
- [2] J. Berntsson and A. Tonderski, "Simple adversarial music generation", <https://github.com/jcberntsson/musicgeneration-project>, 2018.